


Hardware Image Processing

ECE532 Final Presentation
Group 6



1?
6? 2?
2?

The Team

Frank



Cyrus



Savo



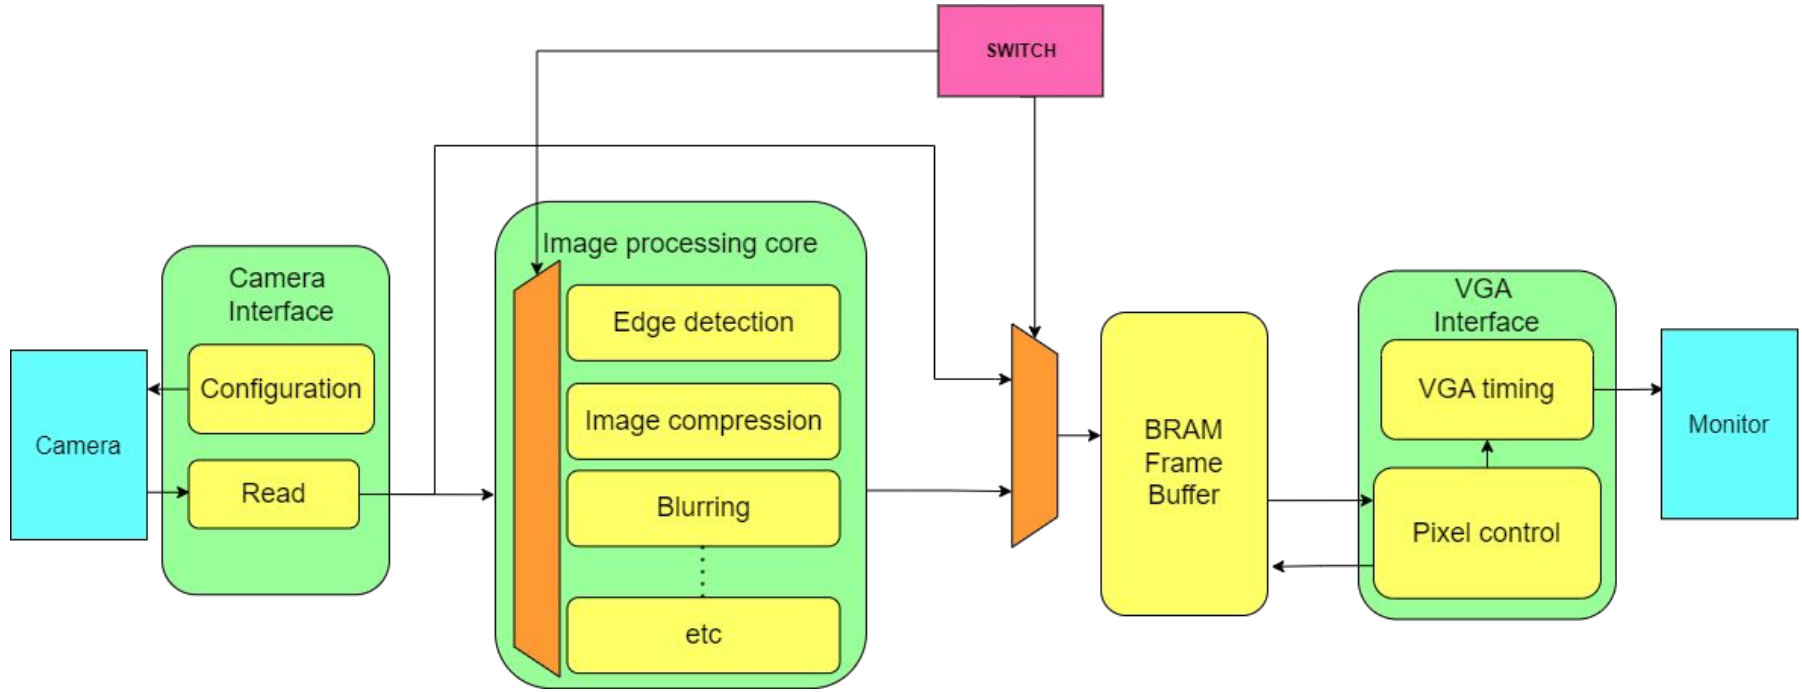
Richard



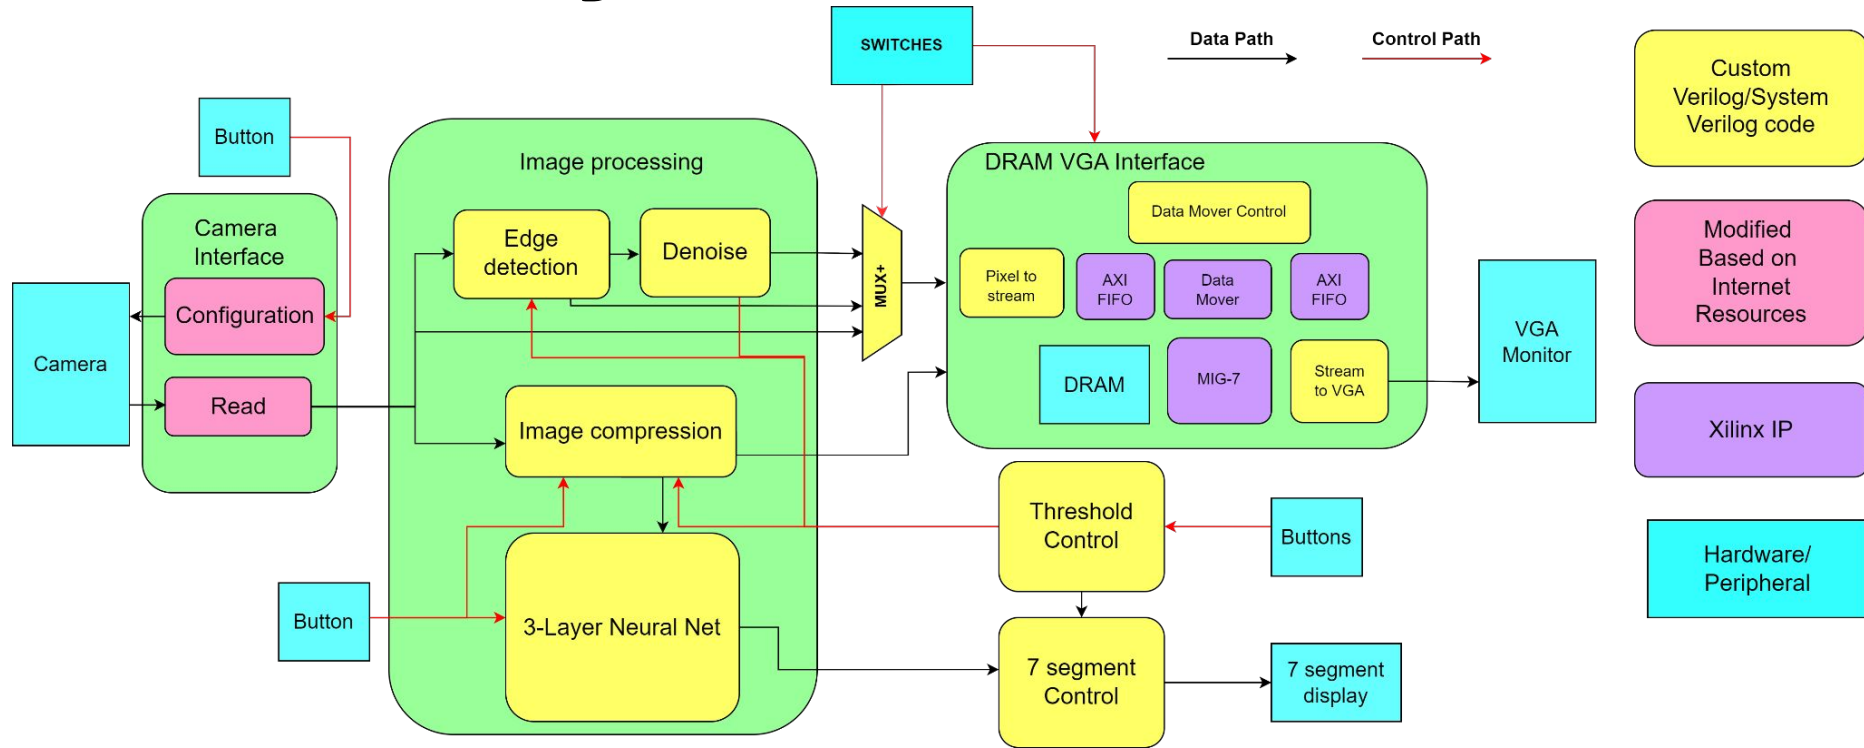
Initial Goals vs Final Project

- Reliable image processing pipeline
 - Changeable kernels
 - Customizable image processing settings
- Reliable image processing pipeline
 - Changeable kernels
 - Customizable image processing settings
 - Digit recognition based on neural network
 - Purely hardware !

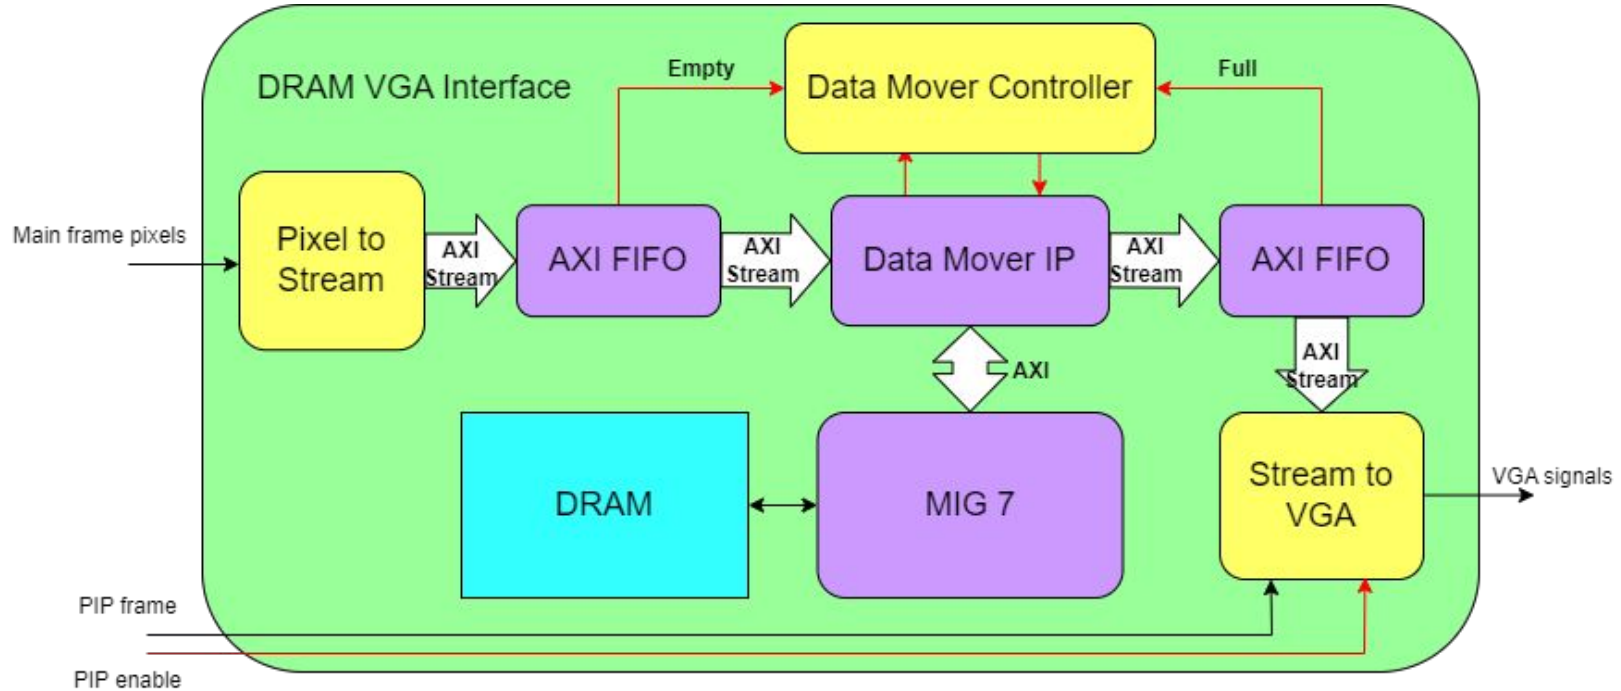
Original Block Diagram (Midproject)



Final Block Diagram



Final Block Diagram - DRAM VGA Interface



Challenges we faced

Data Mover IP Bugs

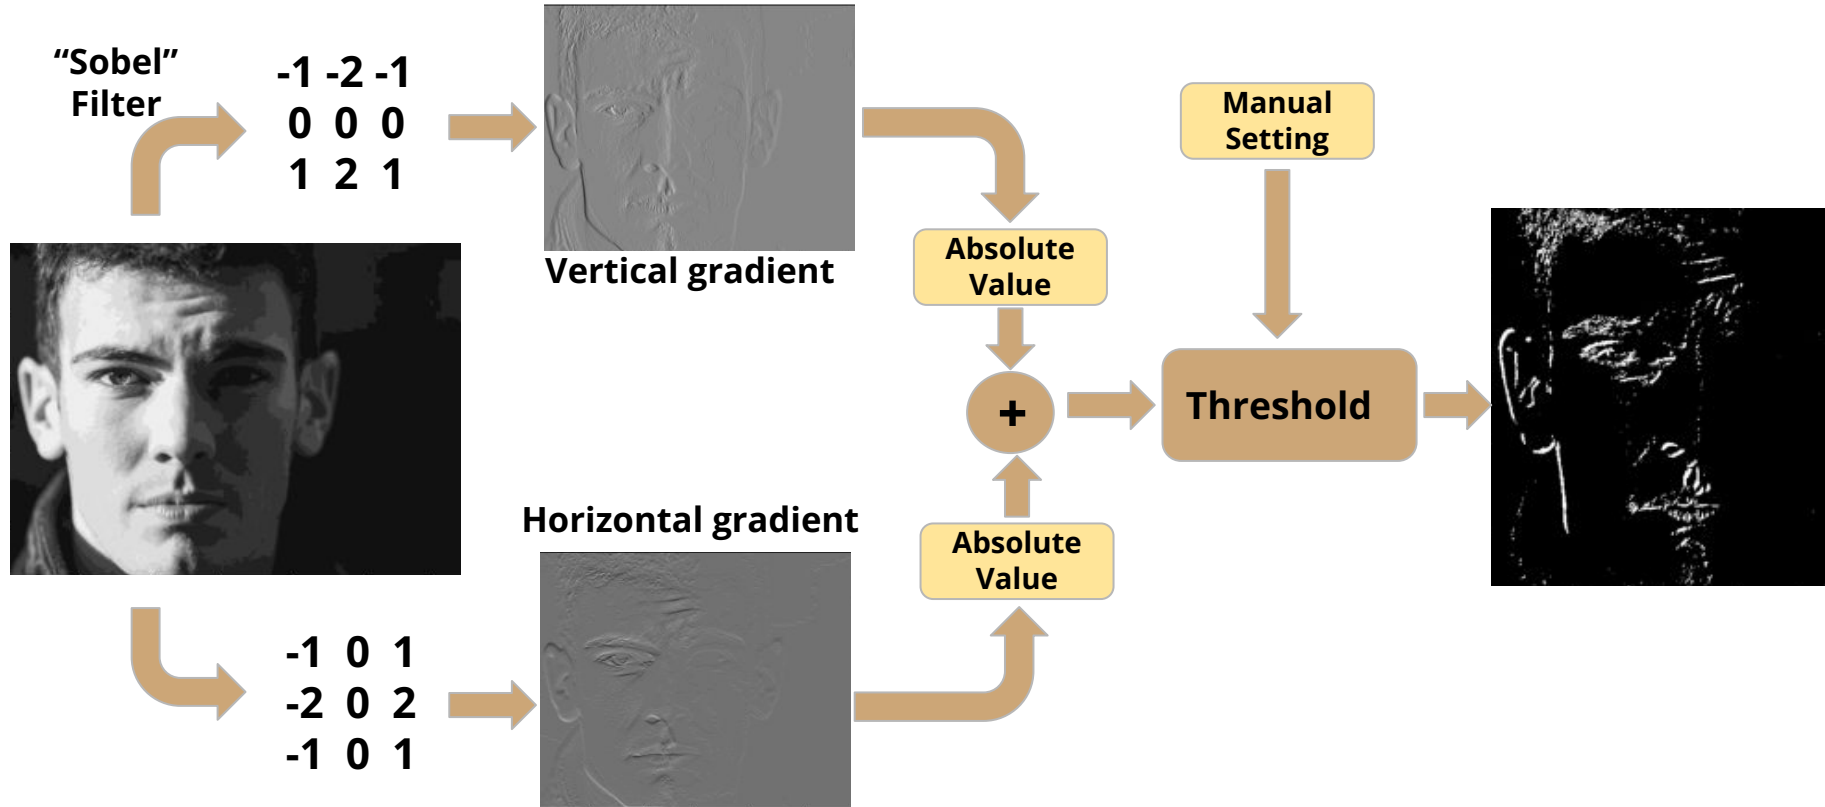
- Not writing data to a continuous region
- Internal buffer on S2MM for two stream packets

0x80000000	0FFF0FFF	0FFF0FFF	AF45C3C3	6697AEE7
0x80000010	0FFF0FFF	0FFF0FFF	65840005	2441AA0C
0X80000020	0FFF0FFF	0FFF0FFF	AD487D2A	080EF016

Accuracy of ML Techniques

- Training Dataset was not taken from our camera + FPGA compression algorithm
- 86% accuracy on test dataset in python does NOT mean 86% accuracy on our system
 - Highly sensitive to position and orientation of number
- More complex Model required to see reasonable accuracy on FPGA

In Depth System Discussion- Edge Detection



In Depth System Discussion- Denoise

The Algorithm:

For each Pixel in the image:


- Check 8 neighboring pixels

- If at least THRESHOLD # neighboring pixels are edges:


 - Accept this pixel as an edge

- Else:

 - This pixel is "noise", set to black



In Depth System Discussion- Image Compression



The Algorithm:

For each 17px square in square crop:

Calculate sum of (15 - pixel intensity) for all pixels in the square

Subtract minimum of sums out of all sums

Scale all sums by a factor so that the maximum is 15

This is compressed pixel intensity

For all pixel intensities:



If pixel compressed pixel intensity < THRESHOLD

Set pixel intensity to 0 (black)

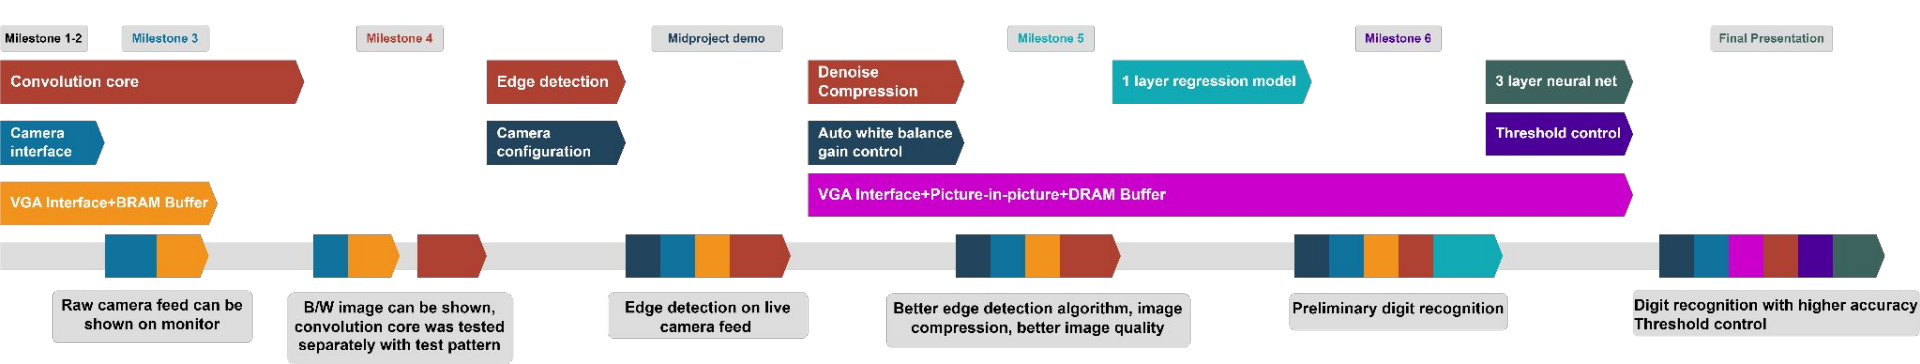
In Depth System Discussion- 3 Layers Neural Net

$(1 \times 784) \times (784 \times 10) \times (10 \times 10) \times (10 \times 10)$

- Linear neural network with biasing
- Trained using gradient descent with backpropagation
- 91.4% accuracy
- 2 hidden layers, 1 output layer



Design Process–A step by step approach



The Final Product

Resource Utilization

- LUTs: 52%
 - Mostly as Logic
 - Also as distributed RAM for pixel buffers
- Flip Flops: 18%
- DSP: 30%
 - Used for Matrix Algebra in Neural Net Classification
- BRAM: 2%
 - Almost No BRAM used; frames stored in DRAM off-chip.

Criterion	Goal	
Throughput	10+ FPS	✓ Full 30 FPS
Latency	<200 ms	✓ ~33ms
Reliability	No dropped connections No pixel glitches	✓

Entirely Hardware, no Microblaze!

Key Takeaways

- (AXI) Streams are powerful
- Properly abiding by interface standards speeds up integration time
- There are many ways to optimize a digital system depending on application
 - DSP vs LUT math
 - Distributed memory, BRAM, DRAM
- Using block diagrams and hierarchies make integration easier
 - Add RTL directly into block diagrams, and wrap systemverilog in verilog code to integrate systemverilog in block diagrams



Demo

Controls

Threshold values for
Edge detection
Denoise
Compression

Leds indicating the
value being changed

Toggle switch for
picture in picture
feature


Predicted
number
(Likelihood):
3rd | 2nd | 1st

Camera
Configuration

Buttons for changing
threshold
Up: Increment by 1
Centre: Change modes
Down: Decrement by 1

Image Capturing
Button

Toggle switches for
Edge detection
Denoise



Q&A

